

```
In [3]: #EXPERIMENT 1
#Implement and demonstrate the FIND-S algorithm for finding the most specific hy
#Read the training data from a csv file.
import pandas as pd
import numpy as np
df=pd.read_csv(r"C:\ml datasets\ml1.csv")
df
concepts=np.array(df)[:,:-1]
concepts
target=np.array(df)[:,-1]
target
def train(c,t):
    for i,val in enumerate(t):
        if val=='yes':
            specific_hypothesis=c[i].copy()
            break
    for i,val in enumerate(c):
        if t[i]=='yes':
            for x in range (len(specific_hypothesis)):
                if val[x]!=specific_hypothesis[x]:
                    specific_hypothesis[x]='?'
    return specific_hypothesis
print("\nThe maximally specific find-S hypothesis for the given training example
print("final hypothesis",train(concepts, target))
```

The maximally specific find-S hypothesis for the given training examples is:  
final hypothesis ['Sunny' 'Warm' '?' 'Strong' '?' '?']


```
In [4]: # EXPERIMENT 2
# implement simple linear regression
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error,r2_score
data=pd.read_csv(r"C:\ml datasets\ml2.csv")
print(data.head())
x=data[['Years of Experience']]
y=data['Salary']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=4)
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
mse=mean_squared_error(y_test,y_pred)
r2=r2_score(y_test,y_pred)
print(f"Mean Squared error:{mse}")
print(f"R Squared error:{r2}")
plt.figure(figsize=(8,6))
plt.scatter(x,y,color='blue',label='Actual Data')
plt.plot(x,model.predict(x),color='red',linewidth=2,label='Regression line')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.title('Linear Regression:Salary vs Years of experience')
plt.legend()
plt.show()
plt.figure(figsize=(8,6))
plt.scatter(x,y,color='blue',label='Actual data')
plt.scatter(x,y,color='red',label='Predicted data')
```

```
plt.xlabel('years')
plt.ylabel('salary')
plt.title("Perfomance:Predict VS Actual Salary")
plt.legend()
plt.show()
```

	Years of Experience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

Mean Squared error:24942746.410306115  
R Squared error:0.9504404484884267





```
In [7]: # EXPERIMENT 3
# implement multi linear regression
import pandas as pd
import matplotlib.pyplot as plt
data=pd.read_csv(r"C:\MLdatasets\ml3.csv")
data.head()
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error,r2_score
x=data[['age','experience']]
y=data['income']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
mse=mean_squared_error(y_test,y_pred)
r2=r2_score(y_test,y_pred)
coefficients = model.coef_
intercept = model.intercept_
print("Mean Squared Error(MSE):",mse)
print ("R2 Score:",r2)
print("Coefficients:",coefficients)
print("Intercept:",intercept)
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.scatter(y_test, y_pred, color='blue')
plt.title('Actual vs Predicted Income')
plt.xlabel('Actual Income')
plt.ylabel('Predicted Income')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', li
plt.subplot(1, 2, 2)
residuals = y_test - y_pred
```

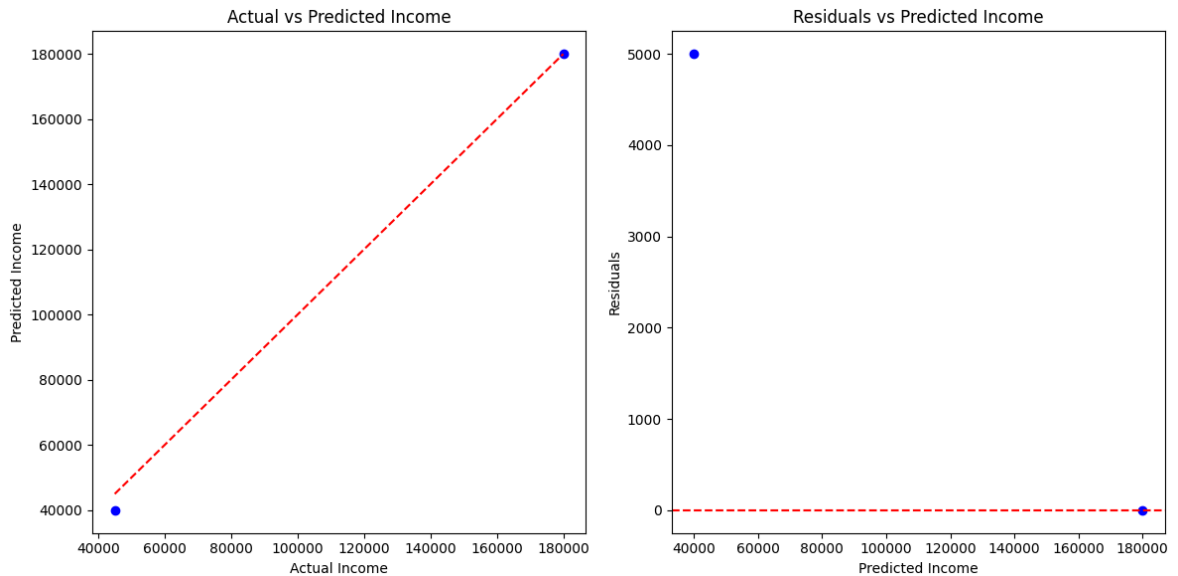
```
plt.scatter(y_pred, residuals, color='blue')
plt.title('Residuals vs Predicted Income')
plt.xlabel('Predicted Income')
plt.ylabel('Residuals')
plt.axhline(y=0, color='red', linestyle='--')
plt.tight_layout()
plt.show()
```

Mean Squared Error(MSE): 12500000.000000145

R2 Score: 0.9972565157750343

Coefficients: [-6000. 10000.]

Intercept: 169999.99999999994



```
In [10]: # EXPERIMENT 4
# Implement Logistic Regression
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score, roc_curve, auc, cl
df = pd.read_csv(r"C:\MLdatasets\ml4.csv")
print(df.head())
print("Class Distribution:\n", df['target'].value_counts())
x = df.drop(columns='target')
y = df['target']
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_
st_x = StandardScaler()
x_train = st_x.fit_transform(x_train)
x_test = st_x.transform(x_test)
classifier = LogisticRegression(class_weight='balanced')
classifier.fit(x_train, y_train)
y_pred = classifier.predict(x_test)
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Pred: No', 'Pre
plt.title('Confusion Matrix')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.show()
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy * 100, "%")
```

```

classreport = classification_report(y_test, y_pred, zero_division=1)
print("\nClassification Report:\n", classreport)
y_prob = classifier.predict_proba(x_test)[: , 1]
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()

```


	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

Class Distribution:

target	
1	5
0	5


Name: count, dtype: int64



Accuracy: 33.33333333333333 %

Classification Report:

	precision	recall	f1-score	support
0	0.50	0.50	0.50	2
1	0.00	0.00	0.00	1
accuracy			0.33	3
macro avg	0.25	0.25	0.25	3
weighted avg	0.33	0.33	0.33	3



```
In [45]: # EXPERIMENT 5
# DATA PREPROCESSING FOR CLASSIFICATION
import numpy as np
import pandas as pd
df=pd.read_csv(r"C:\datasets\l5.csv")
df
x=df.iloc[:, :-1]
x=df.iloc[:, :-1]
x
y=df.iloc[:, -1]
df['AGE']=df["AGE"].fillna(df['AGE'].mean())
df['INCOME']=df['INCOME'].fillna(df['INCOME'].mean())
df
```


Out[45]:

	AGE	INCOME	EDUCATION	GENDER	TARGET
0	32.10	52000.0	HIGH-SCHOOL	MALE	0
1	28.50	48000.0	BACHELOR	FEMALE	1
2	35.00	55000.0	MASTER	MALE	0
3	31.35	53000.0	PHD	FEMALE	1
4	29.80	52000.0	BACHELOR	MALE	0

```
In [16]: # EXPERIMENT 6
# CONFUSION MATRIX FOR BINARY CLASSIFIER
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
from sklearn.preprocessing import StandardScaler
df = pd.read_csv(r"C:\datasets\l6.csv")
x = df.drop(['target'], axis=1)
y = df['target']
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_
scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)
lgr = LogisticRegression(max_iter=1000)
model = lgr.fit(x_train_scaled, y_train)
prediction = model.predict(x_test_scaled)
cm = confusion_matrix(y_test, prediction)
print("Confusion Matrix:")
print(cm)
sns.heatmap(cm, annot=True, cmap='BuPu')
plt.show()
accuracy = accuracy_score(y_test, prediction)
print("Accuracy:", accuracy)
print("Classification Report:")
print(classification_report(y_test, prediction))
```

Confusion Matrix:

```
[[1 1]
 [0 1]]
```



Accuracy: 0.6666666666666666

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.50	0.67	2
1	0.50	1.00	0.67	1
accuracy			0.67	3
macro avg	0.75	0.75	0.67	3
weighted avg	0.83	0.67	0.67	3

```
In [17]: # EXPERIMENT 7
# IMPLEMENT SUPPORT VECTOR MACHINE
import numpy as np
import pandas as pd
from sklearn import datasets
cancer = datasets.load_breast_cancer()
print(cancer.target)
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(cancer.data,cancer.target,test_si
from sklearn import svm
clf = svm.SVC(kernel='linear')
clf.fit(x_train,y_train)
y_pred = clf.predict(x_test)
from sklearn import metrics
print("Accuracy:",metrics.accuracy_score(y_test,y_pred))
print("Precision:",metrics.precision_score(y_test,y_pred))
print("Recall:",metrics.recall_score(y_test,y_pred))
```





```

conf_matrix = confusion_matrix(y_test, y_pred)
ConfusionMatrixDisplay(conf_matrix, display_labels=model.classes_).plot()
plt.show()
print(classification_report(y_test, y_pred, zero_division=0))

```


Initial Data Types and First Few Rows:

```

age          int64
ed           object
income       int64
savings      int64
loan_amount  int64
dtype: object

```

	age	ed	income	savings	loan_amount
0	25	Bachelors	50000	10000	20000
1	30	Masters	60000	15000	25000
2	22	Bachelors	40000	8000	18000
3	28	PhD	70000	20000	30000
4	35	Bachelors	80000	80000	25000

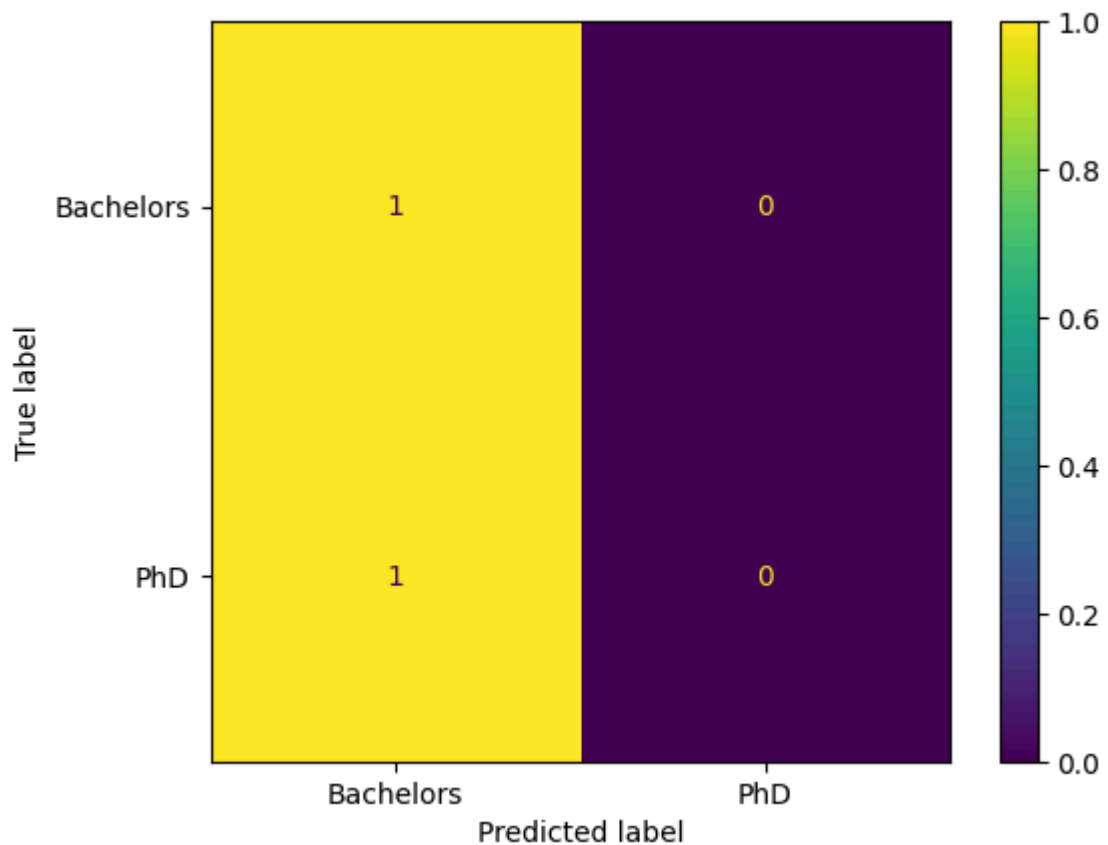


Data after Encoding:

	ed	income	savings	loan_amount	age_25	age_28	age_30	age_35
0	Bachelors	50000	10000	20000	True	False	False	False
1	Masters	60000	15000	25000	False	False	True	False
2	Bachelors	40000	8000	18000	False	False	False	False
3	PhD	70000	20000	30000	False	True	False	False
4	Bachelors	80000	80000	25000	False	False	False	True

Accuracy: 0.5

F1\_score: 0.6666666666666666

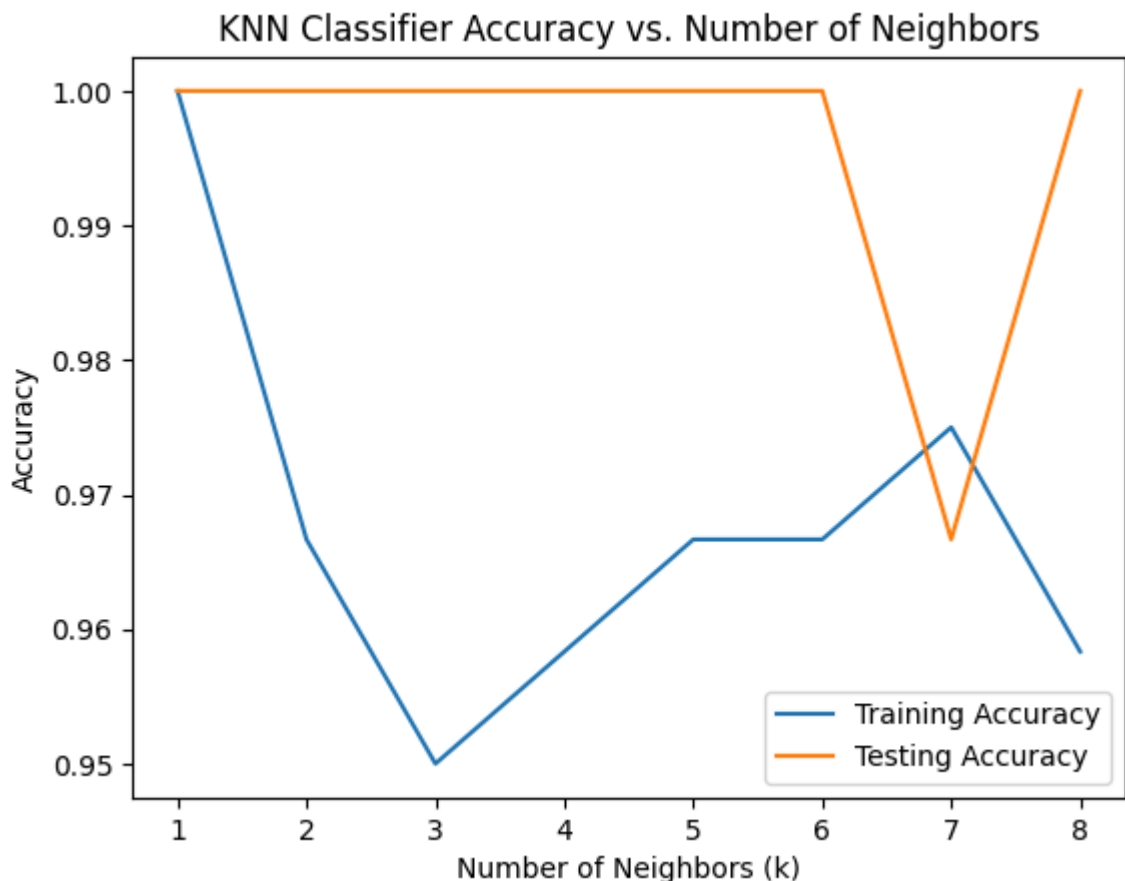


	precision	recall	f1-score	support
Bachelors	0.50	1.00	0.67	1
Masters	0.00	0.00	0.00	1
accuracy				0.50
macro avg	0.25	0.50	0.33	2
weighted avg	0.25	0.50	0.33	2

```
In [31]: # EXPERIMENT 9
# WRITE A PROGRAM TO IMPLEMENT KNEAREST NEIGHBOUR ALGORITHM TO CLASSIFY THE IRIS
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.datasets import load_iris
iris = load_iris()
x = iris.data
y = iris.target
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_
neighbors = np.arange(1, 9)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))
for i, k in enumerate(neighbors):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(x_train, y_train)
    train_accuracy[i] = knn.score(x_train, y_train)
    test_accuracy[i] = knn.score(x_test, y_test)
    y_pred = knn.predict(x_test)
    correct_preds = np.sum(y_pred == y_test)
    incorrect_preds = np.sum(y_pred != y_test)
    print(f"K={k}: Correct Predictions = {correct_preds}, Incorrect Predictions
plt.plot(neighbors, train_accuracy, label='Training Accuracy')
```

```
plt.plot(neighbors, test_accuracy, label='Testing Accuracy')
plt.legend()
plt.xlabel('Number of Neighbors (k)')
plt.ylabel('Accuracy')
plt.title('KNN Classifier Accuracy vs. Number of Neighbors')
plt.show()
```

K=1: Correct Predictions = 30, Incorrect Predictions = 0  
 K=2: Correct Predictions = 30, Incorrect Predictions = 0  
 K=3: Correct Predictions = 30, Incorrect Predictions = 0  
 K=4: Correct Predictions = 30, Incorrect Predictions = 0  
 K=5: Correct Predictions = 30, Incorrect Predictions = 0  
 K=6: Correct Predictions = 30, Incorrect Predictions = 0  
 K=7: Correct Predictions = 29, Incorrect Predictions = 1  
 K=8: Correct Predictions = 30, Incorrect Predictions = 0



```
In [44]: ## EXPERIMENT 10
## WRITE A PROGRAM TO DEMONSTRATE THE WORKING OF THE DECISION TREE BASED ID3 AL
## AN APPROPRIATE DATA THIS KNOWLEDGE TO CLASSIFY A NEW SAMPLE.
import pandas as pd
from sklearn . model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier, export_text , plot_tree
import matplotlib.pyplot as plt
df=pd.read_csv(r"C:\datasets\l10.csv")
df
print("Intial Data Types and First Few Rows:")
print(df.dtypes)
print(df.head())
label_encoders = {}
for column in df.columns:
    if df[column].dtype == 'object' :
        print(f"Encoding column :{column}")
```

```

le = LabelEncoder()
df[column] = le.fit_transform(df[column])
label_encoders[column] = le
print("\n Data after Encoding:")
print(df.head())
x=df.iloc[:, :-1]
y=df.iloc[:, -1]
x
y
x_train,x_test,y_train,y_test =train_test_split(x,y,test_size=0.3, random_state=
id3_tree = DecisionTreeClassifier(criterion="entropy",random_state=42)
id3_tree.fit(x_train,y_train)
accuracy = id3_tree.score (x_test,y_test)
print(f"accuracy :{accuracy:.2f}")
tree_rules = export_text(id3_tree,feature_names=df.columns[: -1].tolist())
print(tree_rules)
import matplotlib.pyplot as plt
from sklearn.tree import plot_tree
plt.figure(figsize=(15, 10))
plot_tree(
    id3_tree,
    feature_names=df.columns[: -1],
    filled=True,
    rounded=True,
    fontsize=12,
    precision=2
)
plt.title("Decision Tree Visualization", fontsize=16)
plt.show()

```

Initial Data Types and First Few Rows:

```

Feature1    object
Feature2    object
Feature3    int64
Target      object
dtype: object

```

	Feature1	Feature2	Feature3	Target
0	A	X	1	Yes
1	B	Y	2	No
2	A	X	1	Yes
3	C	Z	3	No
4	B	Y	2	Yes

Encoding column :Feature1

Encoding column :Feature2

Encoding column :Target

Data after Encoding:

	Feature1	Feature2	Feature3	Target
0	0	0	1	1
1	1	1	2	0
2	0	0	1	1
3	2	2	3	0
4	1	1	2	1

accuracy :0.50


|--- Feature3 <= 2.50

| |--- class: 1

|--- Feature3 > 2.50

| |--- class: 0

## Decision Tree Visualization



In [ ]: